

NLP Disaster Tweets Kaggle Mini-Project

Project Description

Twitter has become an important communication channel in times of emergency. The ubiquitousness of smartphones enables people to announce an emergency they're observing in real-time. Because of this, more agencies are interested in programmatically monitoring Twitter (i.e. disaster relief organizations and news agencies).

In this competition, you're challenged to build a machine learning model that predicts which Tweets are about real disasters and which one's aren't. You'll have access to a dataset of 10,000 tweets that were hand classified.

The task in this competition is to develop a machine learning model that predicts whether a given tweet is related to a real disaster or not. Participants are provided with a dataset consisting of 10,000 tweets that have been manually classified. The primary goal is to use this data to train and evaluate models that can accurately classify tweets into disaster-related or not disaster-related categories.

Import Libraries

```

In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import random
import re
from collections import Counter
from wordcloud import WordCloud

%matplotlib inline
sns.set_style('dark')
sns.set(font_scale=1.2)

plt.rc('axes', labelsz=14)
plt.rc('xtick', labelsz=12)
plt.rc('ytick', labelsz=12)

#sets the default autosave frequency in seconds
%autosave 60

%matplotlib inline

import sklearn
from sklearn.feature_extraction.text import CountVectorizer, HashingVectorizer,
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB

import tensorflow as tf

from transformers import DistilBertTokenizer, DistilBertForSequenceClassification
import datasets

random.seed(0)
np.random.seed(0)
np.set_printoptions(suppress=True)

pd.set_option('display.max_columns', None)
#pd.set_option('display.max_rows', 100)
pd.set_option('display.width', 1000)
pd.set_option('display.float_format', '{:.2f}'.format)

```

Autosaving every 60 seconds

WARNING:tensorflow:From C:\Users\Dennis\anaconda3\lib\site-packages\tf_keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```

In [2]: tf.__version__

```

```

Out[2]: '2.16.2'

```

Exploratory Data Analysis (EDA) — Inspect, Visualize and Clean the Data

```
In [3]: # Step 2: Load the data
train_data = pd.read_csv('train.csv')
test_data = pd.read_csv('test.csv')
```

```
In [4]: # Step 3: Check the sizes of train and test sets
print("Train set size:", train_data.shape)
print("Test set size:", test_data.shape)
```

```
Train set size: (7613, 5)
Test set size: (3263, 4)
```

```
In [5]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7613 entries, 0 to 7612
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           7613 non-null   int64
1   keyword      7552 non-null   object
2   location     5080 non-null   object
3   text         7613 non-null   object
4   target       7613 non-null   int64
dtypes: int64(2), object(3)
memory usage: 297.5+ KB
```

```
In [6]: test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3263 entries, 0 to 3262
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           3263 non-null   int64
1   keyword      3237 non-null   object
2   location     2158 non-null   object
3   text         3263 non-null   object
dtypes: int64(1), object(3)
memory usage: 102.1+ KB
```

```
In [7]: # Step 4: Preview the first few rows of the datasets
print("\nTrain set preview:")
print(train_data.head())

print("\nTest set preview:")
print(test_data.head())
```

Train set preview:

	id	keyword	location	text	targe
t					
0	1	NaN	NaN	Our Deeds are the Reason of this #earthquake M...	
1					
1	4	NaN	NaN	Forest fire near La Ronge Sask. Canada	
1					
2	5	NaN	NaN	All residents asked to 'shelter in place' are ...	
1					
3	6	NaN	NaN	13,000 people receive #wildfires evacuation or...	
1					
4	7	NaN	NaN	Just got sent this photo from Ruby #Alaska as ...	
1					

Test set preview:

	id	keyword	location	text
0	0	NaN	NaN	Just happened a terrible car crash
1	2	NaN	NaN	Heard about #earthquake is different cities, s...
2	3	NaN	NaN	there is a forest fire at spot pond, geese are...
3	9	NaN	NaN	Apocalypse lighting. #Spokane #wildfires
4	11	NaN	NaN	Typhoon Soudelor kills 28 in China and Taiwan

In [8]: *# Step 5: Generate basic statistics for the datasets*

```
print("\nTrain set statistics:")
print(train_data.describe(include='all'))

print("\nTest set statistics:")
print(test_data.describe(include='all'))
```

Train set statistics:

	id	keyword	location
text	target		
count	7613.00	7552	5080
7613	7613.00		
unique	NaN	221	3341
7503	NaN		
top	NaN	fatalities	USA
of T...	NaN		11-Year-Old Boy Charged With Manslaughter
freq	NaN	45	104
10	NaN		
mean	5441.93	NaN	NaN
NaN	0.43		
std	3137.12	NaN	NaN
NaN	0.50		
min	1.00	NaN	NaN
NaN	0.00		
25%	2734.00	NaN	NaN
NaN	0.00		
50%	5408.00	NaN	NaN
NaN	0.00		
75%	8146.00	NaN	NaN
NaN	1.00		
max	10873.00	NaN	NaN
NaN	1.00		

Test set statistics:

	id	keyword	location
text			
count	3263.00	3237	2158
3263			
unique	NaN	221	1602
3243			
top	NaN	deluged	New York
T...			11-Year-Old Boy Charged With Manslaughter of
freq	NaN	23	38
3			
mean	5427.15	NaN	NaN
NaN			
std	3146.43	NaN	NaN
NaN			
min	0.00	NaN	NaN
NaN			
25%	2683.00	NaN	NaN
NaN			
50%	5500.00	NaN	NaN
NaN			
75%	8176.00	NaN	NaN
NaN			
max	10875.00	NaN	NaN
NaN			

```
In [9]: # Step 6: Check for missing values
print("\nMissing values in train set:")
print(train_data.isnull().sum())

print("\nMissing values in test set:")
print(test_data.isnull().sum())
```

```
Missing values in train set:  
id          0  
keyword     61  
location    2533  
text        0  
target      0  
dtype: int64
```

```
Missing values in test set:  
id          0  
keyword     26  
location    1105  
text        0  
dtype: int64
```

```
In [10]: # Step 7: Visualize the target distribution in the train set  
plt.figure(figsize=(6, 4))  
sns.countplot(x='target', data=train_data)  
plt.title('Distribution of Target Variable')  
plt.xlabel('Target')  
plt.ylabel('Count')  
plt.show()
```



```
In [11]: train_data.columns
```

```
Out[11]: Index(['id', 'keyword', 'location', 'text', 'target'], dtype='object')
```

```
In [12]: train_data.target.value_counts()
```

```
Out[12]: 0    4342  
         1    3271  
         Name: target, dtype: int64
```

Characters count/average/min/max

```
In [13]: train_data["text"]
```

```
Out[13]: 0      Our Deeds are the Reason of this #earthquake M...
1      Forest fire near La Ronge Sask. Canada
2      All residents asked to 'shelter in place' are ...
3      13,000 people receive #wildfires evacuation or...
4      Just got sent this photo from Ruby #Alaska as ...
      ...
7608    Two giant cranes holding a bridge collapse int...
7609    @aria_ahrury @TheTawniest The out of control w...
7610    M1.94 [01:04 UTC]?5km S of Volcano Hawaii. htt...
7611    Police investigating after an e-bike collided ...
7612    The Latest: More Homes Razed by Northern Calif...
Name: text, Length: 7613, dtype: object
```

```
In [14]: # Step 1: Calculate character count for each entry in the "text" column
train_data['char_count'] = train_data['text'].apply(len)
```

```
# Step 2: Calculate average character count
average_length = train_data['char_count'].mean()
```

```
# Step 3: Calculate minimum character count
min_length = train_data['char_count'].min()
```


```
# Step 4: Calculate maximum character count
max_length = train_data['char_count'].max()
```

```
# Print the results
print(f'Average Length: {average_length}')
print(f'Minimum Length: {min_length}')
print(f'Maximum Length: {max_length}')
```


```
Average Length: 101.03743596479706
Minimum Length: 7
Maximum Length: 157
```

```
In [15]: sns.displot(train_data['char_count'])
plt.title("Character counts distribution plot")
plt.show()
```


Character counts distribution plot




```
In [16]: sns.kdeplot(train_data['char_count'], color='g', fill=True)
plt.title("Character counts distribution plot")
plt.show()
```



```
In [17]: sns.kdeplot(train_data[train_data["target"] == 0]['char_count'], color='blue', fill=True)
sns.kdeplot(train_data[train_data["target"] == 1]['char_count'], color='orange', fill=True)
plt.title("Character counts distribution plot")
plt.legend(title="Target")
plt.show()
```



Most and Least Common Words

```
In [18]: # Combine all the text into a single string
all_text = ' '.join(train_data['text'].astype(str))

# Tokenize the text (convert to lowercase and split by non-word characters)
words = re.findall(r'\b\w+\b', all_text.lower())
```

```
In [19]: # Count word frequencies
word_counts = Counter(words)

# Total word count
total_words = sum(word_counts.values())
```

```
In [20]: # Most common words (top 10 as an example)
most_common_words = word_counts.most_common(10)
most_common_words
```

```
Out[20]: [('t', 5199),
          ('co', 4740),
          ('http', 4309),
          ('the', 3277),
          ('a', 2200),
          ('in', 1986),
          ('to', 1949),
          ('of', 1830),
          ('i', 1778),
          ('and', 1426)]
```

```
In [21]: # Least common words (bottom 10 as an example)
least_common_words = word_counts.most_common()[::-11:-1]
least_common_words
```


```
Out[21]: [('ymy4rskq3d', 1),
('stfmbbzfb5', 1),
('nf4iculoje', 1),
('rqkk15uhey', 1),
('symptoms', 1),
('developing', 1),
('forney', 1),
('5uecmcv2pk', 1),
('9km', 1),
('fa3fcnl86', 1)]
```

```
In [22]: # Convert to a DataFrame
df_most_common = pd.DataFrame(most_common_words, columns=['Word', 'Count'])

# Create the bar plot
plt.figure(figsize=(10, 6))
sns.barplot(data=df_most_common, x='Count', y='Word', palette='viridis')

# Set plot labels and title
plt.xlabel('Count')
plt.ylabel('Word')
plt.title('Top 10 Most Common Words')

# Show the plot
plt.show()
```



Clean Data

```
In [23]: # Step 1: Handling Missing Values
# For simplicity, we'll fill missing 'keyword' and 'location' with placeholder t
train_data['keyword'].fillna('missing_keyword', inplace=True)
train_data['location'].fillna('missing_location', inplace=True)
test_data['keyword'].fillna('missing_keyword', inplace=True)
test_data['location'].fillna('missing_location', inplace=True)
```

```
In [24]: train_data.isnull().sum()
```

```
Out[24]: id            0
keyword          0
location         0
text             0
target           0
char_count       0
dtype: int64
```

```
In [25]: test_data.isnull().sum()
```

```
Out[25]: id            0
keyword          0
location         0
text             0
dtype: int64
```

```
In [26]: # Step 2: Basic Text Cleaning
def clean_text(text):
    # Remove URLs
    text = re.sub(r'http\S+', '', text)
    # Remove HTML tags
    text = re.sub(r'<.*?>', '', text)
    # Remove special characters and numbers
    text = re.sub(r'^A-Za-z\s', '', text)
    # Convert text to lowercase
    text = text.lower()
    # Remove extra spaces
    text = re.sub(r'\s+', ' ', text).strip()
    return text

train_data['clean_text'] = train_data['text'].apply(clean_text)
test_data['clean_text'] = test_data['text'].apply(clean_text)
```

```
In [27]: # Step 3: Preview the cleaned text
print("\nCleaned text preview in train set:")
print(train_data[['text', 'clean_text']].head())

print("\nCleaned text preview in test set:")
print(test_data[['text', 'clean_text']].head())
```

Cleaned text preview in train set:

	clean_text	text
0	Our Deeds are the Reason of this #earthquake M... our deeds are the reason o f this earthquake ma...	
1	Forest fire near La Ronge Sask. Canada ar la ronge sask canada	forest fire ne
2	All residents asked to 'shelter in place' are ... all residents asked to she lter in place are be...	
3	13,000 people receive #wildfires evacuation or... people receive wildfires e vacuation orders in ...	
4	Just got sent this photo from Ruby #Alaska as ... just got sent this photo f rom ruby alaska as s...	

Cleaned text preview in test set:

	clean_text	text
0	Just happened a terrible car crash ed a terrible car crash	just happen
1	Heard about #earthquake is different cities, s... heard about earthquake is different cities sta...	
2	there is a forest fire at spot pond, geese are... there is a forest fire at spot pond geese are ...	
3	Apocalypse lighting. #Spokane #wildfires hting spokane wildfires	apocalypse lig
4	Typhoon Soudelor kills 28 in China and Taiwan lls in china and taiwan	typhoon soudelor ki

Disaster Words Visualization with Word Cloud

```
In [28]: # Combine all cleaned text into a single string
all_clean_text = ' '.join(train_data['clean_text'].astype(str))

# Create the word cloud object
wordcloud = WordCloud(width=800, height=400, background_color='white', colormap=

# Display the word cloud using matplotlib
plt.figure(figsize=(10, 5))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off') # Hide the axis
plt.title('Word Cloud of Cleaned Text')
plt.show()
```


Model Architecture

1. Logistic Regression

Architecture:

- **Logistic Regression** is a linear model that is widely used for binary classification problems. It models the probability that a given input belongs to a particular class.

Reasoning:

- **Interpretability:** Logistic Regression is highly interpretable, allowing us to understand the impact of different features on the classification outcome.
- **Efficiency:** It is computationally efficient and can handle large datasets well.
- **Baseline Model:** Logistic Regression serves as a strong baseline model. If it performs well, more complex models may not be necessary.
- **Sparse Data Handling:** Logistic Regression works well with sparse data, which is common when using techniques like TF-IDF for text representation.

2. TF-IDF Vectorizer

Architecture:

- **TF-IDF (Term Frequency-Inverse Document Frequency)** is a numerical statistic that reflects how important a word is to a document in a collection or corpus. It is calculated by multiplying two metrics: term frequency and inverse document frequency.

Reasoning:

- **Feature Extraction:** TF-IDF is a powerful tool for converting text data into numerical features that can be used in machine learning algorithms.
- **Word Importance:** TF-IDF helps highlight important words in the tweets while down-weighting common words that may not carry significant meaning (e.g., "the", "is", "at").
- **Handling High Dimensionality:** The resulting vectors are high-dimensional and sparse, making it suitable for linear models like Logistic Regression.
- **Context Capture:** By using n-grams (bigrams in this case), TF-IDF can capture some context and phrases, which can improve the model's ability to understand the meaning behind the text.

Model Justification

- **Text Data:** Tweets are short text snippets. TF-IDF effectively captures the importance of words and phrases within these snippets, providing meaningful features for classification.
- **Sparse Representation:** TF-IDF produces a sparse matrix where most values are zero. Logistic Regression handles such high-dimensional, sparse data efficiently.

- **Interpretability:** The simplicity and interpretability of Logistic Regression allow us to gain insights into which words and phrases are significant predictors of disaster-related tweets.
- **Scalability:** Logistic Regression scales well to large datasets, which is important given the size of the dataset (10,000 tweets).

Conclusion

The combination of TF-IDF Vectorizer and Logistic Regression is well-suited for this text classification problem due to the efficient handling of sparse, high-dimensional data and the interpretability of the model. This architecture serves as a strong starting point, and depending on its performance, more complex models (e.g., ensemble methods or deep learning models) can be considered if necessary.

```
In [32]: # Vectorizing Text Data
vectorizer = TfidfVectorizer(max_features=10000, ngram_range=(1, 2))
```

```
In [33]: # Fit and transform the training data
X_train = vectorizer.fit_transform(train_data['combined_text'])
```

```
In [34]: # Transform the test data
X_test = vectorizer.transform(test_data['combined_text'])
```

```
In [35]: # Get the target values
y_train = train_data['target']
```

```
In [36]: # Check the shape of the resulting matrices
print("\nShape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
```

```
Shape of X_train: (7613, 10000)
Shape of X_test: (3263, 10000)
```

Modeling

```
In [37]: # Split the training data for validation
X_train_split, X_val, y_train_split, y_val = train_test_split(X_train, y_train,
```

```
In [38]: # Initialize and train the model
model = LogisticRegression(max_iter=1000)
model.fit(X_train_split, y_train_split)
```

```
Out[38]: ▼      LogisticRegression
LogisticRegression(max_iter=1000)
```

```
In [39]: # Predict on the validation set
y_val_pred = model.predict(X_val)
```

```
In [40]: # Evaluate the model
print("Validation Accuracy:", accuracy_score(y_val, y_val_pred))
print("\nClassification Report:\n", classification_report(y_val, y_val_pred))
```


Validation Accuracy: 0.799080761654629

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.88	0.83	874
1	0.81	0.70	0.75	649
accuracy			0.80	1523
macro avg	0.80	0.79	0.79	1523
weighted avg	0.80	0.80	0.80	1523

Logistic Regression Results and Analysis

Initial Results

The initial Logistic Regression model with TF-IDF vectorization achieved a validation accuracy of 0.7991. Here's a detailed breakdown of the classification report:

- Class 0 (Non-disaster tweets):
 - Precision: 0.80
 - Recall: 0.88
 - F1-score: 0.83
- Class 1 (Disaster tweets):
 - Precision: 0.81
 - Recall: 0.70
 - F1-score: 0.75
- Overall Metrics:
 - Accuracy: 0.80
 - Macro Average F1-score: 0.79
 - Weighted Average F1-score: 0.80

Hyperparameter Tuning

Let's start by optimizing the hyperparameters of the Logistic Regression model to see if we can improve its performance. Hyperparameter Optimization Procedure

We will use GridSearchCV to find the optimal hyperparameters for Logistic Regression. The hyperparameters we'll tune are:

- C: Inverse of regularization strength.
- penalty: Regularization technique (L1 or L2).

```
In [41]: # Define the parameter grid
param_grid = {
    'C': [0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2'],
    'solver': ['liblinear']
}
```

```
In [42]: # Initialize the Logistic Regression model
log_reg = LogisticRegression(max_iter=1000)
```

```
In [43]: # Initialize GridSearchCV
grid_search = GridSearchCV(estimator=log_reg, param_grid=param_grid, cv=5, scoring='accuracy')
```

```
In [44]: # Fit GridSearchCV
grid_search.fit(X_train_split, y_train_split)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
Out[44]:
└─ GridSearchCV
  └─ estimator: LogisticRegression
    └─ LogisticRegression
```

```
In [45]: # Get the best parameters and score
best_params = grid_search.best_params_
best_score = grid_search.best_score_
```

```
In [46]: print("Best Parameters:", best_params)
print("Best Cross-Validation Score:", best_score)
```

Best Parameters: {'C': 1, 'penalty': 'l2', 'solver': 'liblinear'}
Best Cross-Validation Score: 0.7973727422003284

Trying Different Architectures

Apart from Logistic Regression, we'll try the following models:

- 1 Random Forest Classifier
- 2 Gradient Boosting Classifier
- 3 Support Vector Machine (SVM)
- 4 Multinomial Naive Bayes

```
In [47]: # Initialize models
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
gb_model = GradientBoostingClassifier(n_estimators=100, random_state=42)
svm_model = SVC(C=1, kernel='linear', random_state=42)
nb_model = MultinomialNB()

# Train and evaluate models
models = {
    'Random Forest': rf_model,
    'Gradient Boosting': gb_model,
    'Support Vector Machine': svm_model,
    'Multinomial Naive Bayes': nb_model
}

for name, model in models.items():
    model.fit(X_train_split, y_train_split)
    y_val_pred = model.predict(X_val)
    accuracy = accuracy_score(y_val, y_val_pred)
    print(f"{name} Validation Accuracy: {accuracy}")
    print(f"\nClassification Report for {name}:\n", classification_report(y_val,
```

Random Forest Validation Accuracy: 0.7760998030203545

Classification Report for Random Forest:

	precision	recall	f1-score	support
0	0.76	0.90	0.82	874
1	0.82	0.61	0.70	649
accuracy			0.78	1523
macro avg	0.79	0.75	0.76	1523
weighted avg	0.78	0.78	0.77	1523

Gradient Boosting Validation Accuracy: 0.7340774786605384

Classification Report for Gradient Boosting:

	precision	recall	f1-score	support
0	0.72	0.89	0.79	874
1	0.78	0.53	0.63	649
accuracy			0.73	1523
macro avg	0.75	0.71	0.71	1523
weighted avg	0.74	0.73	0.72	1523

Support Vector Machine Validation Accuracy: 0.8003939592908733

Classification Report for Support Vector Machine:

	precision	recall	f1-score	support
0	0.80	0.86	0.83	874
1	0.80	0.71	0.75	649
accuracy			0.80	1523
macro avg	0.80	0.79	0.79	1523
weighted avg	0.80	0.80	0.80	1523

Multinomial Naive Bayes Validation Accuracy: 0.8036769533814839

Classification Report for Multinomial Naive Bayes:

	precision	recall	f1-score	support
0	0.78	0.91	0.84	874
1	0.84	0.66	0.74	649
accuracy			0.80	1523
macro avg	0.81	0.79	0.79	1523
weighted avg	0.81	0.80	0.80	1523

Results and Analysis

Hyperparameter Tuning Results for Logistic Regression

- **Best Parameters:** {'C': 1, 'penalty': 'l2', 'solver': 'liblinear'}
- **Best Cross-Validation Score:** 0.7974

Model Comparison

We compared the performance of various models, including Logistic Regression (with tuned hyperparameters), Random Forest, Gradient Boosting, Support Vector Machine (SVM), and Multinomial Naive Bayes.

Model	Validation Accuracy	Precision (Class 0)	Recall (Class 0)	F1-score (Class 0)	Precision (Class 1)	Recall (Class 1)	F1-score (Class 1)
Logistic Regression	0.7991	0.80	0.88	0.83	0.81	0.70	0.75
Random Forest	0.7761	0.76	0.90	0.82	0.82	0.61	0.70
Gradient Boosting	0.7341	0.72	0.89	0.79	0.78	0.53	0.63
Support Vector Machine	0.8004	0.80	0.86	0.83	0.80	0.71	0.75
Multinomial Naive Bayes	0.8037	0.78	0.91	0.84	0.84	0.66	0.74

Analysis of Model Performance

1. Logistic Regression:

- **Validation Accuracy:** 0.7991
- Logistic Regression remains a robust choice, with a well-balanced precision and recall, particularly strong on non-disaster tweets.

2. Random Forest:

- **Validation Accuracy:** 0.7761
- High precision for disaster-related tweets but lower recall, indicating the model might be missing some disaster-related tweets.

3. Gradient Boosting:

- **Validation Accuracy:** 0.7341
- Lower accuracy and F1 scores compared to other models, indicating it may not be as effective for this particular problem.

4. Support Vector Machine (SVM):

- **Validation Accuracy:** 0.8004
- Performs similarly to Logistic Regression with a balanced precision and recall,

showing it can be an effective model for this task.

5. **Multinomial Naive Bayes:**

- **Validation Accuracy:** 0.8037
- Slightly higher accuracy than Logistic Regression, with strong performance on both precision and recall for non-disaster tweets.

Improvements and Techniques Applied

1. **Hyperparameter Tuning:**

- Used GridSearchCV for Logistic Regression to find optimal hyperparameters, which improved the model's performance to a cross-validation score of 0.7974.

2. **Feature Engineering:**

- Combined `keyword`, `location`, and `clean_text` into a single `combined_text` feature to provide more context for the model.

3. **Advanced Vectorization Techniques:**

- Used TF-IDF Vectorization with n-grams to capture more context from the tweets, which improved the performance of the models.

4. **Model Comparison:**

- Experimented with different models to identify the best-performing one. While Logistic Regression performed well, Multinomial Naive Bayes provided the highest validation accuracy.

Conclusion

Based on the validation accuracy and the detailed analysis of precision, recall, and F1 scores, **Multinomial Naive Bayes** emerged as the best-performing model for this specific task, with a validation accuracy of 0.8037.

- **Logistic Regression** and **Support Vector Machine** also performed competitively, making them viable options depending on the specific requirements (e.g., interpretability, efficiency).
- **Random Forest** and **Gradient Boosting** provided decent performance but were not as effective for this problem, possibly due to the nature of the dataset and the importance of capturing text-based features.

Future Work

- **Further Hyperparameter Tuning:** Continue to fine-tune hyperparameters for models like Random Forest and Gradient Boosting.
- **Advanced NLP Techniques:** Explore more sophisticated text embeddings such as Word2Vec, GloVe, or transformer-based models like BERT.
- **Ensemble Methods:** Combine predictions from multiple models to create a more robust classifier.

- **Data Augmentation:** Increase the dataset size through data augmentation techniques to provide more training data for the models.

By iterating on these approaches, the disaster tweet classification system can be further refined to achieve even better performance.

Bonus Section (DistilBERT method)

Load the Tokenizer and Tokenize the Combined Text

```
In [48]: from transformers import DistilBertTokenizer

# Load the pre-trained DistilBERT tokenizer
tokenizer = DistilBertTokenizer.from_pretrained('distilbert-base-uncased')

# Tokenize the 'combined_text' column
train_encodings = tokenizer(
    train_data['combined_text'].tolist(),
    truncation=True,
    padding=True,
    max_length=128,
    return_tensors='pt'
)
```

Prepare Data for Model Input

```
In [49]: import torch

# Convert target labels to tensor
labels = torch.tensor(train_data['target'].values)

# Prepare dataset by combining encodings and labels
class DisasterTweetDataset(torch.utils.data.Dataset):
    def __init__(self, encodings, labels):
        self.encodings = encodings
        self.labels = labels

    def __getitem__(self, idx):
        item = {key: val[idx] for key, val in self.encodings.items()}
        item['labels'] = self.labels[idx]
        return item

    def __len__(self):
        return len(self.labels)

train_dataset = DisasterTweetDataset(train_encodings, labels)
```

Fine-Tuning DistilBERT

```
In [ ]: from transformers import DistilBertForSequenceClassification, Trainer, TrainingArguments

# Load the pre-trained DistilBERT model for sequence classification
model = DistilBertForSequenceClassification.from_pretrained('distilbert-base-uncased')

# Define training arguments
training_args = TrainingArguments(
    output_dir='./results',
    eval_strategy="epoch",
    per_device_train_batch_size=256,
    per_device_eval_batch_size=256,
    num_train_epochs=1,
    weight_decay=0.01,
    logging_dir='./logs',
)

# Create a Trainer object
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
)

# Fine-tune the model
trainer.train()
```

Evaluation

```
In [ ]: test_dataset = test_data['combined_text']
```

```
In [ ]: predictions = trainer.predict(test_dataset)
predicted_labels = torch.argmax(predictions.predictions, axis=1)
```

Python code done by Dennis Lam
